

AWS Labs – S3, Sagemaker, Devops.

- Amazon S3:
 - Scalable object storage service used for storing files, data backups, and hosting static websites.
- Amazon SageMaker:
 - A cloud service for building, training, and deploying machine learning models at scale.
- DevOps in AWS:
 - Tools like AWS CodePipeline, CodeBuild, and CodeDeploy streamline CI/CD processes and infrastructure automation.

Demo:

Amazon S3 and CloudFront: Static Website Deployment

Task 1: Create an S3 Bucket for Static Website Hosting

Step 1: Log into the AWS Management Console

- Go to [AWS Console](#).
- Type "S3" in the search bar, and open Amazon S3.

Step 2: Create a New S3 Bucket

- Click Create bucket.
- Enter a unique Bucket name (e.g., **user10-static-website-demo**).
- Select a Region closest to your target audience, such as Mumbai (ap-south-1).
- Under Block Public Access settings, uncheck Block all public access to allow public website access. Confirm the acknowledgment.
- Click Create bucket to finish.

Step 3: Upload Website Files

- Download the HTML and CSS files from [this source link](#).
- Extract the files on your computer.
- In your S3 bucket, click Upload and Add files.
- Select the HTML and other files and Upload them.
- Rename the main HTML file to index.html (Actions > Rename Object).

Step 4: Enable Static Website Hosting

- Go to the Properties tab of the bucket.
- Scroll to Static website hosting and click Edit.
- Select Enable.
- Set Index Document as `index.html`.
- (Optional) Enter an error document (e.g., `404.html`) if you have one.
- Click Save changes.

Step 5: Make Files Public

- Go to the Permissions tab.
- Scroll to Bucket Policy and click Edit.
- Paste the following bucket policy, replacing `<YOUR-BUCKET-NAME-HERE>` with your bucket name:

```
{  
  
  "Version": "2012-10-17",  
  
  "Statement": [  
  
    {  
  
      "Sid": "granting-access-to-s3-objects",  
  
      "Principal": "*",  
  
      "Effect": "Allow",  
  
      "Action": ["s3:GetObject"],  
  
      "Resource": "arn:aws:s3:::<YOUR-BUCKET-NAME-HERE>/*"  
  
    }  
  
  ]  
  
}
```

Click Save changes.

Step 6: Test the Static Website

- Copy the Endpoint URL from the Static website hosting section.
- Open the URL in a browser to confirm that your website is accessible.

Task 2: Create a CloudFront Distribution for the S3 Bucket

Step 1: Navigate to CloudFront

- In the AWS Console, search for CloudFront and open it.
- Click Create Distribution.

Step 2: Configure the Origin Settings

- For Origin Domain, select your S3 bucket.
- Leave Origin Path blank unless files are stored in a subfolder.
- Under Origin access, select Create new OAC to create Origin Access Control, and update the S3 bucket policy to allow CloudFront access.
- Under Default cache behavior, set CachingOptimized for the cache policy.
- In Settings, select Use all edge locations for maximum performance.
- Set Default root object to `index.html`.
- Click Create Distribution. It may take a few minutes to deploy.

Task 3: Update S3 Bucket Permissions for CloudFront Access Only

- In CloudFront, observe the yellow notification to update the bucket policy for OAC.
- Copy the suggested policy and open S3.
- Go to your bucket's Permissions > Bucket Policy > Edit.
- Replace the existing policy with the copied OAC policy and Save changes.

Task 4: Test the CloudFront Distribution

- Wait until the CloudFront distribution status is Deployed.
- Copy the CloudFront domain name (e.g., `xxxxxxxxxx.cloudfront.net`).
- Open this URL in a browser to check if the website is accessible via CloudFront.
- Refresh to see the caching effect from CloudFront.

Task 5: Delete Resources to Avoid Costs

- Delete CloudFront Distribution:
 - Disable the CloudFront distribution.
 - Wait until the status changes to Disabled, then delete it.
- Delete the S3 Bucket:
 - Empty the bucket by deleting all files, then delete the bucket itself.

Amazon SageMaker: Iris Dataset Classification

Step 1: Prepare the Dataset

Use Python to load, split, and save the dataset:

```
from sklearn.datasets import load_iris
import pandas as pd
from sklearn.model_selection import train_test_split

iris = load_iris()
data = pd.DataFrame(data=iris['data'],
                    columns=iris['feature_names'])
data['target'] = iris['target']

train, test = train_test_split(data, test_size=0.2,
                               random_state=42)

train.to_csv('train.csv', index=False)
test.to_csv('test.csv', index=False)
```

Step 2: Upload Data to S3

Open the AWS CLI or AWS Console and create an S3 bucket if you don't already have one:

```
aws s3 mb s3://your-bucket-name
```

Upload the files:

```
aws s3 cp train.csv s3://your-bucket-name/train/train.csv
```

```
aws s3 cp test.csv s3://your-bucket-name/test/test.csv
```

Step 3: Create a Model in SageMaker

In SageMaker, create an XGBoost Estimator for model training:

```
from sagemaker.estimator import Estimator
import boto3
import sagemaker

role = 'your-sagemaker-role'
xgboost_container = sagemaker.image_uris.retrieve("xgboost",
boto3.Session().region_name, "1.2-1")

estimator = Estimator(
    image_uri=xgboost_container,
    role=role,
    instance_count=1,
    instance_type='ml.m5.large',
    output_path='s3://your-bucket/output/',
    hyperparameters={
        'max_depth': 3,
        'eta': 0.2,
        'objective': 'multi:softmax',
        'num_class': 3,
        'num_round': 100,
    }
)
```

Step 4: Train the Model

```
s3_input_train = 's3://your-bucket-name/train/train.csv'
estimator.fit({'train': s3_input_train})
```

Step 5: Deploy the Model

```
predictor = estimator.deploy(  
    initial_instance_count=1,  
    instance_type='ml.m5.large'  
)
```

Step 6: Test the Model

Send sample data to the deployed model:

```
iris_sample = [5.1, 3.5, 1.4, 0.2] # Example data  
result = predictor.predict(iris_sample)  
print(result) # Predicted class
```

DevOps: Deploy Web Application Using AWS DevOps Tools

Step 1: Fork the Repository

Log into GitHub, navigate to the [aws-elastic-beanstalk-express-js-sample](#) repository, and select Fork.

Step 2: Deploy the Web App to Elastic Beanstalk

- In the AWS Elastic Beanstalk Console, select Create Application.
- Choose a Web server environment and enter an Application Name (e.g., **DevOpsGettingStarted**).
- Set the Platform to Node.js, and confirm other settings as default.
- Review the application and environment configuration, and then submit to create.

Step 3: Create a Build Project in CodeBuild

- In the AWS CodeBuild Console, select Create project.
- Name the project (e.g., **Build-DevOpsGettingStarted**) and select GitHub as the Source Provider.
- Authenticate with GitHub and choose the repository you forked.
- Configure the environment with Amazon Linux 2 and Runtime: Standard. Use `aws/codebuild/amazonlinux2-x86_64-standard:3.0` as the image.

Use the following `buildspec.yml` file:

```
version: 0.2
phases:
  build:
    commands:
      - npm i --save
artifacts:
  files:
    - '**/*'
```

Step 4: Set Up a Delivery Pipeline in CodePipeline

- In the AWS CodePipeline Console, choose Create pipeline.
- Select GitHub as the Source provider and authenticate to access the forked repository.
- Configure AWS CodeBuild as the Build provider and choose your `Build-DevOpsGettingStarted` project.
- For the Deploy provider, select AWS Elastic Beanstalk, and use `DevOpsGettingStarted-env` for the environment.
- Review and create the pipeline.

Step 5: Test the Pipeline

- After creation, watch the pipeline run and confirm each step completes successfully.
- Once deployed, open AWS Elastic Beanstalk, go to your environment, and click on the URL to see the deployed web app.

RESULTS - Screenshots

S3:

AWS Services: Amazon S3 and Amazon CloudFront

Amazon S3 - Simple Storage Service



AWS Sagemaker:

The screenshot displays the AWS SageMaker console interface. On the left, a sidebar shows navigation options: Applications (6), Home, Running instances, Data, Auto ML, Experiments, Jobs, Pipelines, Models, JumpStart, and Deployments. The main area is titled "JupyterLab" and includes a "Create JupyterLab space" button. Below the title, there's an "About" section and a search bar. A table lists the running JupyterLab spaces:

Name	Application	Status	Type	Last modified	Action
crimson-jls-01	JupyterLab	Running	Private	10 days ago	Stop Open

Below the table, it indicates "1 results" and "Results are cached". A "Refresh" button is present. At the bottom, there's a section titled "Introducing spaces (New)" with a "Learn more" link.

File Edit View Run Kernel Git Tabs Settings Help default-20241023T181878 / crimson-jis-01

crimson_model.ipynb x Terminal 1 x +

Notebook Cluster Python 3 (ipykernel)

```
[1]: # Data Preparation
# =====
import sagemaker
import boto3
import pandas as pd
from sklearn.model_selection import train_test_split
import os

# Initialize SageMaker session and specify bucket
sagemaker_session = sagemaker.Session()
bucket = sagemaker_session.default_bucket()
role = sagemaker.get_execution_role()

# Load Iris dataset
iris_data = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data',
header=None)
iris_data.columns = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width', 'species']
# Convert species to numeric values
iris_data['species'] = iris_data['species'].astype('category').cat.codes
print(iris_data['species'].unique()) # Should output [0, 1, 2]
# Split data into train and validation sets
train_data, val_data = train_test_split(iris_data, test_size=0.2, random_state=42)
# Move the label column to the front of the dataframe
train_data = train_data[['species', 'sepal_length', 'sepal_width', 'petal_length', 'petal_width']]
val_data = val_data[['species', 'sepal_length', 'sepal_width', 'petal_length', 'petal_width']]
# Save locally as CSV
train_data.to_csv('train.csv', header=False, index=False)
val_data.to_csv('validation.csv', header=False, index=False)
# One can print the train_data and the val_data that would be the 80% training data 20% validation data

sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location: /home/sagemaker-user/.config/sagemaker/config.yaml
[0 1 2]

[2]: # upload the train_data and val_data to S3 bucket
s3_train_path = sagemaker_session.upload_data(path='train.csv', bucket=bucket, key_prefix='xgboost-iris/train')
s3_validation_path = sagemaker_session.upload_data(path='validation.csv', bucket=bucket, key_prefix='xgboost-iris/validation')
print(f"Training data uploaded to: {s3_train_path}")
print(f"Validation data uploaded to: {s3_validation_path}")
```

Simple 1 1 42% Amazon Q Cookie Preferences 1

File Edit View Run Kernel Git Tabs Settings Help default-20241023T181878 / crimson-jis-01

crimson_model.ipynb x Terminal 1 x +

Notebook Cluster Python 3 (ipykernel)

```
# Prepare a sample input for prediction
test_data = val_data.drop(columns=['species']).iloc[2].values # Assuming 'species' is the label column
test_data = ','.join(map(str, test_data)) # Convert features to a CSV row format
print(f"Test data for prediction: {test_data}")

# Make a prediction
response = xgboost_predictor.predict(
test_data,
initial_args={'ContentType': 'text/csv'}) # Specify content type as text/csv

# The response might need decoding, depending on format
predicted_class = response.decode('utf-8') # Decode the response to a readable format (if required)
print(f"Predicted class: {predicted_class}")

# Clean up by deleting the endpoint
xgboost_predictor.delete_endpoint()

INFO:sagemaker:Creating model with name: sagemaker-xgboost-2024-10-24-07-17-16-600
INFO:sagemaker:Creating endpoint-config with name sagemaker-xgboost-2024-10-24-07-17-16-600
INFO:sagemaker:Creating endpoint with name sagemaker-xgboost-2024-10-24-07-17-16-600
-----!Test data for prediction: 7.7,2.6,6.9,2.3
Predicted class: 2.0

INFO:sagemaker:Deleting endpoint configuration with name: sagemaker-xgboost-2024-10-24-07-17-16-600
Predicted class: 2.0

INFO:sagemaker:Deleting endpoint with name: sagemaker-xgboost-2024-10-24-07-17-16-600

Click to add a cell.
```

Simple 1 1 42% Amazon Q Cookie Preferences 1